

MindsOnML



<https://www.linkedin.com/in/desaimann37/>




<https://github.com/desaimann37/MindsOnML>

Hope this helps!

David Desai

MLLM (Multi-modal Large Language Models)



Disadvantages of previous work (CLIP models)

↳ (1)

Separate Encoders for image & text resulting in
"Modality Gap."

[Contrastive Language
Image pre-training]





Figure 5: Distribution of image embeddings and text embeddings with different representation methods.

↳ (2) Cannot handle Interleaved Image/text Inputs.

Interleaved

Input



Instead of this (non-interleaved):

vbnet

[Copy code](#)

```
Text: "A dog is running in the park."  
Image: dog.jpg
```

Interleaved input looks like:

css

[Copy code](#)


```
"The dog starts running" → [IMAGE_1] →  
"then jumps over a log" → [IMAGE_2]
```

The model processes this as one continuous sequence.

↳ ③

Needs Massive Image-text Pairs for training (~ 400 million pairs to train Standard CLIP models)

CLIP training ↴



- Here, both embeddings generated are of same dimension (R^d) but, they are different in values and to make in "same space", we need contrastive loss to calculate loss form $\text{image} \rightarrow \text{text}$ and $\text{text} \rightarrow \text{image}$.

Step ① Compute pairwise similarity.

$$S_{ij} = \text{v-img}_i \cdot \text{v-txt}_j$$

$(N \times N)$

for each $\text{img}(i)$

↳ it's matching text should have highest similarity.

Step ② Loss for $(\text{image} \rightarrow \text{text})$

$$L_{\text{img}} = \text{CrossEntropy}(\text{softmax}(S_i, :), \text{target}=i)$$

- Loss for $(\text{text} \rightarrow \text{image})$

$$L_{\text{text}} = \text{CrossEntropy}(\text{softmax}(S_-, :, i), \text{target}=i)$$

Step ③ Final Loss:

$$L = (L_{\text{img}} + L_{\text{text}}) / 2$$

↳ this forces alignment.

- Till this point, CLIP processes img-text pairs and calculate loss (contrastive) like this and these disadvantages are along with them.
- So, we need to move to MLLMs to overcome these issues.

Simple Intuition of E5V (Using MLLMs):

① Force MLLM to convert any Input
(Image or text or Image + text) into
meaning based embeddings using prompts.

② Use prompts like:

Given Image: Summary of above image in one word.

Given Text : summary of text sentence in one word.

③ Take last token embedding as the representation
This compresses meaning into a single token,
not modality.

④ Why this works: MLLM understands semantic
not pixels or words.

- The prompt removes the "modality gap", hence text & image fall into same semantic space.
- Embeddings align by meaning, not the input type.

⑤ Training strategy: (very important contribution)

- Train only on text pairs (images are optional to give)
- Use contrastive learning on text-text pairs (unlike, CLIP based models where image-text pairs are used)
- During training:
 - ↳ Remove visual encoder.
 - ↳ Train only the LLM.

⑥ What E5V can do?

- Text-image Retrieval
- Image-Image Retrieval
- Composed Image (text + image) Retrieval
- Sentence Embeddings.
- All in zero-shot setting.

⑦ At Inference time:

- ↳ Visual encoder & projector is available
So, if we give image optionally, LLM will convert to text embeddings.

⇒ Results Obtained :

we use this to find cosine similarity between cat & dog.

- text embedding
- image embedding
- interleaved embedding

Inference



Summary above image in one word:
cat

LLM

Projection
Vision Encoder

MLLM



A brown and white dog is standing:
dog

dog

our code does have this image as input.

GPU available

```
CUDA available: True
GPU count: 1
GPU name: NVIDIA A100-SXM4-80GB MIG 3g.40gb
Compute capability: (8, 0)
VRAM (GB): 42.144366592
(gclip) [mdesai23@gpu019 testing]$ python e5v.py
Loading processor...
Loading model...
Loading checkpoint shards: 100% |██████████| 4/4 [00:04<00:00,  1.12s/it]
Could not cache non-existence of file. Will ignore error and continue. Error: [Errno 122] Disk quota exceeded: '/home/mdesai23/.cache/huggingface/hub/models--llm-hf--llama3-llava-next-8b-hf/.no exist/b041c0d0ea0dd0196d147206c210c8d175fc2da/custom_generate'
Model loaded on: cuda:0
Vision encoder + projector frozen.

--- E5-V SANITY CHECK (INFERENCE) ---
Embedding shape: torch.Size([1, 4096])
sim(text dog, text animal): 0.89900234375
sim(text dog, text car): 0.79296875
sim(image dog, text dog): 0.63623046875
sim(image dog, image cat): 0.8779296875
```


Final
Result

LLAVA - NEXT (8B)

{hugging face}

```
E5-V image + text inference working correctly.
(gclip) [mdesai23@gpu019 testing]$
```

```
t1 = embed_text("A dog is running")
t2 = embed_text("An animal is moving fast")
t3 = embed_text("A parked car")
```

Training Data

- 558K filtered image-text pairs from LAION/CC/SBU, captioned by BLIP.
- 158K GPT-generated multimodal instruction-following data.
- 500K academic-task-oriented VQA data mixture.
- 50K GPT-4V data mixture.
- 40K ShareGPT data.

Additionally,

This python file (esv.py) shows only how given two images are similar (using cosine similarity)

- But, we can also use it for,

Contrastive question answering (QA)

```
image_labels = ["dog.jpg", "cat.jpg"]

# Contrastive questions
queries = [
    1. ("an image of a dog", "an image of a cat"),
    2. ("an image of a cat", "an image of a dog"),
    3. ("an image of an animal", "an image of a vehicle"),
    # Semantic logic
    4. ("an image of a pet", "an image of an object"),
    5. ("a living thing", "a non-living object"),
    6. ("a domestic animal", "a wild animal"),
]
```

Output :

① dog.jpg , ② Cat.jpg , ③ dog.jpg , ④ Cat.jpg , ⑤ Cat.jpg , ⑥ dog.jpg

This gives image labels with matched result given by esv2.py

- All python files are attached in linked in post as well as on github Repo.

e5v. py :

```
import torch
import torch.nn.functional as F
from transformers import AutoProcessor, AutoModelForImageTextToText
from PIL import Image
```

```
# =====#
# CONFIG
# =====#
MODEL_ID = "llava-hf/llama3-llava-next-8b-hf"
DTYPE = torch.float16
```

```
# =====#
# PROMPTS (FROM E5-V PAPER)
# =====#
TEXT_PROMPT = """{text}
Summary of the above sentence in one word:"""
IMAGE_PROMPT = """<image>
Summary of the above image in one word:"""
# =====#
```

```
# =====#
# LOAD MODEL
# =====#
print("Loading processor...")
processor = AutoProcessor.from_pretrained(
    MODEL_ID,
    use_fast=True
)
```

```
print("Loading model...")
model = AutoModelForImageTextToText.from_pretrained(
    MODEL_ID,
    dtype=DTYPE,
    device_map="auto"
)
```

```
model.eval()
print("Model loaded on:", next(model.parameters()).device)
```

```
# =====#
# FREEZE VISION ENCODER (PAPER REQUIREMENT)
# =====#
```

```
if hasattr(model, "vision_tower"):
    for p in model.vision_tower.parameters():
        p.requires_grad = False
```

```
if hasattr(model, "multi_modal_projector"):
    for p in model.multi_modal_projector.parameters():
        p.requires_grad = False
```

```
print("Vision encoder + projector frozen.")
```

```
# =====#
# EMBEDDING FUNCTIONS (INFERENCE ONLY)
# =====#
```

```
@torch.no_grad()
def embed_text(text: str) -> torch.Tensor:
    prompt = TEXT_PROMPT.format(text=text)

    inputs = processor(
        text=prompt,
        return_tensors="pt"
    ).to(model.device)
```

```
outputs = model(
    *inputs,
    output_hidden_states=True,
    return_dict=True
)
```

```
emb = outputs.hidden_states[-1][:, -1, :]
emb = F.normalize(emb, dim=-1)
return emb
```

```
@torch.no_grad()
def embed_image(image: Image.Image) -> torch.Tensor:
    inputs = processor(
        images=image,
        text=IMAGE_PROMPT,
        return_tensors="pt"
    ).to(model.device)
```

```
outputs = model(
    *inputs,
    output_hidden_states=True,
    return_dict=True
)
```

```
emb = outputs.hidden_states[-1][:, -1, :]
emb = F.normalize(emb, dim=-1)
return emb
```

```
# =====#
# SANITY CHECK (TEXT + IMAGE)
# =====#
if __name__ == "__main__":
    print("\n--- E5-V SANITY CHECK (INFERENCE) ---")
```

```
# ----- TEXT à TEXT -----
t1 = embed_text("A dog is running")
t2 = embed_text("An animal is moving fast")
t3 = embed_text("A parked car")
```

```
print("Embedding shape:", t1.shape)
print("sim(text dog, text animal):", F.cosine_similarity(t1, t2).item())
print("sim(text dog, text car):", F.cosine_similarity(t1, t3).item())
```

```
# ----- IMAGE à TEXT -----
# Replace with your own image path
img = Image.open("/home/mdesai23/project/MLLM/testing/cat_test.jpg").convert("RGB")
```

```
img_emb = embed_image(img)
txt_emb = embed_text("a dog running")
```

```
print("sim(image dog, text dog):",
      F.cosine_similarity(img_emb, txt_emb).item())
```

```
# ----- IMAGE à IMAGE (OPTIONAL) -----
img2 = Image.open("/home/mdesai23/project/MLLM/testing/dog_test.jpg").convert("RGB")
```

```
img_emb2 = embed_image(img2)
print("sim(image dog, image cat):",
      F.cosine_similarity(img_emb, img_emb2).item())
```

```
print("\nE5-V image + text inference working correctly.")
```

```

e5v2.py: import torch
         import torch.nn.functional as F
         from transformers import AutoProcessor, AutoModelForImageTextToText

         MODEL_ID = "llava-hf/llama3-llava-next-8b-hf"
         DTYPE = torch.float16

# =====#
# PROMPTS (FROM E5-V PAPER)
# =====#

         TEXT_PROMPT = """{text}
Summary of the above sentence in one word:"""

         IMAGE_PROMPT = """<image>
Summary of the above image in one word:"""

# =====#
# LOAD MODEL
# =====#

         print("Loading processor...")
         processor = AutoProcessor.from_pretrained(
             MODEL_ID,
             use_fast=True
         )

         print("Loading model...")
         model = AutoModelForImageTextToText.from_pretrained(
             MODEL_ID,
             dtype=DTYPE,
             device_map="auto"
         )

         model.eval()
         print("Model loaded on:", next(model.parameters()).device)

# =====#
# FREEZE VISION ENCODER (PAPER REQUIREMENT)
# =====#

         if hasattr(model, "vision_tower"):
             for p in model.vision_tower.parameters():
                 p.requires_grad = False

         if hasattr(model, "multi_modal_projector"):
             for p in model.multi_modal_projector.parameters():
                 p.requires_grad = False

         print("Vision encoder + projector frozen.")

# =====#
# EMBEDDING FUNCTIONS (INFERENCE ONLY)
# =====#

         @torch.no_grad()
         def embed_text(text: str) -> torch.Tensor:
             prompt = TEXT_PROMPT.format(text=text)

             inputs = processor(
                 text=prompt,
                 return_tensors="pt"
             ).to(model.device)

             outputs = model(
                 **inputs,
                 output_hidden_states=True,
                 return_dict=True
             )

             emb = outputs.hidden_states[-1][:, -1, :]
             return F.normalize(emb, dim=-1)

         @torch.no_grad()
         def embed_image(image: Image.Image) -> torch.Tensor:
             inputs = processor(
                 images=image,
                 text=IMAGE_PROMPT,
                 return_tensors="pt"
             ).to(model.device)

             outputs = model(
                 **inputs,
                 output_hidden_states=True,
                 return_dict=True
             )

             emb = outputs.hidden_states[-1][:, -1, :]
             return F.normalize(emb, dim=-1)

```

```

# =====#
# CONTRASTIVE RETRIEVAL QA (FIXED OPTION A)
# =====#

         # answer_question_contrastive(
         #     positive_query: str,
         #     negative_query: str,
         #     image_embeddings: torch.Tensor,
         #     image_labels: list
         # ):
         #     ...
         #     Contrastive retrieval:
         #     score = sim(pos_query, image) - sim(neg_query, image)
         #     ...

         pos_emb = embed_text(positive_query)
         neg_emb = embed_text(negative_query)

         pos_scores = F.cosine_similarity(pos_emb, image_embeddings)
         neg_scores = F.cosine_similarity(neg_emb, image_embeddings)

         scores = pos_scores - neg_scores
         best_idx = scores.argmax().item()

         return image_labels[best_idx], scores[best_idx].item()

# =====#
# DEMO / SANITY CHECK
# =====#

         if __name__ == "__main__":
             print("\n--- E5-V CONTRASTIVE RETRIEVAL QA DEMO ---")

             # Load images
             img_dog = Image.open(
                 "/home/mdesai23/project/MLLM/testing/dog_test.jpg"
             ).convert("RGB")

             img_cat = Image.open(
                 "/home/mdesai23/project/MLLM/testing/cat_test.jpg"
             ).convert("RGB")

             # Embed images
             emb_dog = embed_image(img_dog)
             emb_cat = embed_image(img_cat)

             image_embeddings = torch.cat([emb_dog, emb_cat], dim=0)
             image_labels = ["dog.jpg", "cat.jpg"]

             # Contrastive questions
             queries = [
                 ("an image of a dog", "an image of a cat"),
                 ("an image of a cat", "an image of a dog"),
                 ("an image of an animal", "an image of a vehicle"),
                 # Semantic logic
                 ("an image of a pet", "an image of an object"),
                 ("a living thing", "a non-living object"),
                 ("a domestic animal", "a wild animal"),
             ]

             for pos_q, neg_q in queries:
                 label, score = answer_question_contrastive(
                     pos_q,
                     neg_q,
                     image_embeddings,
                     image_labels
                 )

                 print(f"Positive query: {pos_q}")
                 print(f"Negative query: {neg_q}")
                 print(f"Answer: {label} (score={score:.3f})\n")

```